

Lab Assignment 4

Convolutional Neural Networks

Eleftherios Karamoulas - S3261859
Panagiotis Tzafos - S3302148

March 26, 2017

2 Theory questions

1. The pooling layer serves to progressively reduce the spatial size of the representation, to reduce the number of parameters and amount of computation in the network, and hence to also control overfitting(CS231n lecture). In practice, its use is to downsample its input, reducing the amount of resources that the network needs in order to perform the computations by reducing the dimensions of the input using a filter, a stride and an elementwise activation function. Also, after the downsampling it is obvious that the network will have less parameters. As a result, it will be able to generalize better in new situations and avoid overfitting(small training set error, large error for new examples).
2. Weight sharing is a very important feature, as it can dramatically reduce the number of weights(less computations). The idea behind this is that the number of unique sets of weights can be equal to the number of filters that are going to be used, as we are about use these to every different region of the image we are scanning. This is when the network is interested to examine if a feature exists in the image indepentantly by its position. It can be applied because of the fact that, assuming that a single depth slice(filter) weight configuration is associated with a single feature that our network is looking for(edges, circles etc.), in every single one spatial region that the filter checks, it responsible of tracing these specific features(Not useful when the network has to learn in different positions of the image different features, e.g. a training set of faces, that are centralized).
3. The time is less in ReLu activation function compared to the sigmoid and the tanh activation functions, as it does not involve expensive operations(exponentials, etc.), as it can be implemented by simply thresholding a matrix of activations at zero. The results are better because of the properties that ReLu function provides, like sparse representation(the more activations are ≤ 0 the more sparse is our model and this is generally more beneficial than dense representations that sigmoid and tanh provide) and ReLu function non-saturating form(sigmoid and tanh functions are vanishing gradient, as the absolute value of x in $\alpha = Wx + b$ increases the gradient of sigmoid becomes increasingly small compared to ReLu that the gradient is constant.)
4. (a) From the given data and the formula to compute the output we have that $W = 12, F = 3, P = 0, S = 1$, so output $= \frac{W-F+2*P}{S} + 1 = \frac{12-3+2*0}{1} + 1 = 10$ and because we have 3 filters the total neurons will be $10 \times 10 \times 3 = 300$. The total number of parameters will be 300×1 weights for the full depth(grayscale) of the input image $+ 3$ for the 3 biases of each filters $= 303$.

(b) The input layer is our image and if we assume that each pixel is one neuron, our input layer consists of $12 \times 12 = 144$ neurons and each of them will be fully connected to

our neurons in the hidden layer which consists of 300 neurons. As a result the total connections between input and hidden layer will be $144 \times 300 = 43200$ weights + 300 biases = 43500. It is obvious that the number of parameters of the fully-connected layer compared with a convolutional layer is much bigger, and that makes convolutional layer more efficient.

5. This approach to solve the car decision problem is questionable because even though we can use CNN's and get a network that can decide between cars an average solution, the features that characterise the car can be subjective for different customers depending on their needs and comforts.

4 Convolutional Layer

```

1 function convolvedFeatures = cnnConvolve(filterDim, numFilters, images, W, b)
2 %cnnConvolve Returns the convolution of the features given by W and b with
3 %the given images
4 %
5 % Parameters:
6 % filterDim - filter (feature) dimension
7 % numFilters - number of feature maps
8 % images - large images to convolve with, matrix in the form
9 %           images(r, c, image number)
10 % W, b - W, b for features from the sparse autoencoder
11 %        W is of shape (filterDim,filterDim,numFilters)
12 %        b is of shape (numFilters,1)
13 %
14 % Returns:
15 % convolvedFeatures - matrix of convolved features in the form
16 %                    convolvedFeatures(imageRow, imageCol, featureNum, imageNum)
17
18 numImages = size(images, 3);
19 imageDim = size(images, 1);
20 convDim = imageDim - filterDim + 1;
21 convolvedFeatures = zeros(convDim, convDim, numFilters, numImages);
22
23 % Instructions:
24 % Convolve every filter with every image here to produce the
25 % (imageDim - filterDim + 1) x (imageDim - filterDim + 1) x numFeatures x numImages
26 % matrix convolvedFeatures, such that
27 % convolvedFeatures(imageRow, imageCol, featureNum, imageNum) is the
28 % value of the convolved featureNum feature for the imageNum image over
29 % the region (imageRow, imageCol) to (imageRow + filterDim - 1, imageCol + filterDim -
30 % 1)
31 %
32 % Expected running times:
33 % Convolving with 100 images should take less than 30 seconds
34 % Convolving with 5000 images should take around 2 minutes
35 % (So to save time when testing, you should convolve with less images, as
36 % described earlier)
37
38 %%% Add code here
39 for i=1:numImages

```

```

39     im = images(:,:,i);
40     for j=1:numFilters
41         cw = rot90(W(:,:,j),2);
42         cf = conv2(im,cw,'valid') + b(j);
43         convolvedFeatures(:,:,j,i) = 1./(1+exp(-cf));
44     end
45 end
46 end

```

Listing 1: cnnConvolve.m

5 Mean Pooling Layer

```

1 function pooledFeatures = cnnPool(poolDim, convolvedFeatures)
2 %cnnPool Pools the given convolved features
3 %
4 % Parameters:
5 % poolDim - dimension of pooling region
6 % convolvedFeatures - convolved features to pool (as given by cnnConvolve)
7 %
8 %               convolvedFeatures(imageRow, imageCol, featureNum, imageNum)
9 % Returns:
10 % pooledFeatures - matrix of pooled features in the form
11 %
12 %               pooledFeatures(poolRow, poolCol, featureNum, imageNum)
13 %
14 numImages = size(convolvedFeatures, 4);
15 numFilters = size(convolvedFeatures, 3);
16 convolvedDim = size(convolvedFeatures, 1);
17 pooledFeatures = zeros(convolvedDim / poolDim, ...
18     convolvedDim / poolDim, numFilters, numImages);
19 % Instructions:
20 % Now pool the convolved features in regions of poolDim x poolDim,
21 % to obtain the
22 % (convolvedDim/poolDim) x (convolvedDim/poolDim) x numFeatures x numImages
23 % matrix pooledFeatures, such that
24 % pooledFeatures(poolRow, poolCol, featureNum, imageNum) is the
25 % value of the featureNum feature for the imageNum image pooled over the
26 % corresponding (poolRow, poolCol) pooling region.
27 %
28 % Use mean pooling here.
29
30 %%% Add code here
31 for i=1:numImages
32     for j=1:numFilters
33         x=1;
34         y=1;
35         for k=1:poolDim:convolvedDim
36             for l=1:poolDim:convolvedDim
37                 pooledFeatures(x,y,j,i) = mean2(convolvedFeatures...
38                     (k:k+poolDim-1,l:l+poolDim-1,j,i));
39                 if(y==(convolvedDim / poolDim))
40                     x=x+1;
41                     y=1;

```

```

42         else
43             y = y+1;
44         end
45     end
46 end
47 end
48 end
49 end

```

Listing 2: cnnPool.m

6 Forward Pass

```

1 function [cost, grad, preds, activations] = cnnCost(theta,images,labels,numClasses,...
2             filterDim,numFilters,poolDim,pred)
3 % Calcualte cost and gradient for a single layer convolutional
4 % neural network followed by a softmax layer with cross entropy
5 % objective.
6 %
7 % Parameters:
8 % theta      - unrolled parameter vector
9 % images     - stores images in imageDim x imageDim x numImges
10 %             array
11 % numClasses - number of classes to predict
12 % filterDim  - dimension of convolutional filter
13 % numFilters - number of convolutional filters
14 % poolDim    - dimension of pooling area
15 % pred       - boolean only forward propagate and return
16 %             predictions
17 %
18 %
19 % Returns:
20 % cost       - cross entropy cost
21 % grad       - gradient with respect to theta (if pred==False)
22 % preds      - list of predictions for each example (if pred==True)
23
24
25 if ~exist('pred','var')
26     pred = false;
27 end;
28
29
30 imageDim = size(images,1); % height/width of image
31 numImages = size(images,3); % number of images
32
33 %% Reshape parameters and setup gradient matrices
34
35 % Wc is filterDim x filterDim x numFilters parameter matrix
36 % bc is the corresponding bias
37
38 % Wd is numClasses x hiddenSize parameter matrix where hiddenSize
39 % is the number of output units from the convolutional layer
40 % bd is corresponding bias
41 [Wc, Wd, bc, bd] = cnnParamsToStack(theta,imageDim,filterDim,numFilters,...

```

```

42         poolDim,numClasses);
43
44 % Same sizes as Wc,Wd,bc,bd. Used to hold gradient w.r.t above params.
45 Wc_grad = zeros(size(Wc));
46 Wd_grad = zeros(size(Wd));
47 bc_grad = zeros(size(bc));
48 bd_grad = zeros(size(bd));
49
50 %=====
51 %% Forward Propagation
52 % In this step you will forward propagate the input through the
53 % convolutional and subsampling (mean pooling) layers. You will then use
54 % the responses from the convolution and pooling layer as the input to a
55 % standard softmax layer.
56
57 %% Convolutional Layer
58 % For each image and each filter, convolve the image with the filter, add
59 % the bias and apply the sigmoid nonlinearity. Then subsample the
60 % convolved activations with mean pooling. Store the results of the
61 % convolution in activations and the results of the pooling in
62 % activationsPooled. You will need to save the convolved activations for
63 % backpropagation.
64 convDim = imageDim-filterDim+1; % dimension of convolved output
65 outputDim = (convDim)/poolDim; % dimension of subsampled output
66
67 % convDim x convDim x numFilters x numImages tensor for storing activations
68
69 %%% REPLACE THE FOLLOWNG LINE %%%
70 activations = cnnConvolve(filterDim, numFilters, images, Wc, bc);
71
72 % outputDim x outputDim x numFilters x numImages tensor for storing
73 % subsampled activations
74
75 %%% REPLACE THE FOLLOWNG LINE %%%
76 activationsPooled = cnnPool(poolDim,activations);
77
78 % Reshape activations into 2-d matrix, hiddenSize x numImages,
79 % for Softmax layer
80 %%% REPLACE THE FOLLOWING LINE %%%
81 activationsPooled = reshape(activationsPooled,[],numImages);
82
83 %% Softmax Layer
84 % Forward propagate the pooled activations calculated above into a
85 % standard softmax layer. For your convenience we have reshaped
86 % activationPooled into a hiddenSize x numImages matrix. Store the
87 % results in probs.
88
89 % numClasses x numImages for storing probability that each image belongs to
90 % each class.
91
92 %%% COMPUTE THE SOFTMAX OUTPUT %%%
93 probs = zeros(numClasses,numImages);
94 Ywx = Wd*activationsPooled;

```

```

95 Ywxb = bsxfun(@plus,Ywx,bd);
96 Ynum = exp(Ywxb);
97 probs = bsxfun(@divide,Ynum,sum(Ynum,1));
98
99 %=====
100 %% STEP 1b: Calculate Cost
101 % In this step you will use the labels given as input and the probs
102 % calculate above to evaluate the cross entropy objective. Store your
103 % results in cost.
104 indexes = sub2ind(size(probs), labels', 1:numImages);
105 cost = -mean(log(probs(indexes)));
106
107 % Makes predictions given probs and returns without backproagating errors.
108 [~,preds] = max(probs,[],1);
109 preds = preds';
110 if pred
111     grad = 0;
112     return;
113 end;
114
115 %=====
116 %% STEP 1c: Backpropagation
117 % Backpropagate errors through the softmax and convolutional/subsampling
118 % layers. Store the errors for the next step to calculate the gradient.
119 % Backpropagating the error w.r.t the softmax layer is as usual. To
120 % backpropagate through the pooling layer, you will need to upsample the
121 % error with respect to the pooling layer for each filter and each image.
122 % Use the kron function and a matrix of ones to do this upsampling
123 % quickly.
124
125 deriv = probs;
126 deriv(indexes) = deriv(indexes) - 1;
127 deriv = deriv ./ numImages;
128
129 Wd_grad = deriv * activationsPooled';
130 bd_grad = sum(deriv, 2);
131
132 deriv2_pooled = Wd' * deriv;
133 deriv2_pooled = reshape(deriv2_pooled, outputDim, outputDim, numFilters, numImages);
134 delta_upsampled = zeros(convDim, convDim, numFilters, numImages);
135
136 for im_idx=1:numImages
137     im = squeeze(images(:,:,im_idx));
138     for f_idx=1:numFilters
139         delta_pool = (1/poolDim^2) * kron(squeeze(deriv2_pooled(:,:,f_idx,im_idx)), ones
            (poolDim));
140         delta_upsampled(:,:,f_idx, im_idx) = delta_pool .* ...
            activations(:,:,f_idx,im_idx).*(1-activations(:,:,f_idx,im_idx));
141         delta_pool_sqz = squeeze(delta_upsampled(:,:,f_idx,im_idx));
142         cur_grad = conv2(im, rot90(delta_pool_sqz, 2), 'valid');
143
144         Wc_grad(:,:,f_idx) = Wc_grad(:,:,f_idx) + cur_grad;
145         bc_grad(f_idx) = bc_grad(f_idx) + sum(delta_pool_sqz(:));
146

```

```

147     end
148 end
149
150 %% Unroll gradient into grad vector for minFunc
151 grad = [Wc_grad(:) ; Wd_grad(:) ; bc_grad(:) ; bd_grad(:)];
152
153 end

```

Listing 3: cnnCost.m

7 Experiments

1. **Question 1**
2. In the beginning of the training we see that our filters look like random noise patches but as our training progresses we see that the filters achieve more structure and finally after 3 epochs we can see clearly on them shapes e.g.(lines, angles, curves) that they can recognize on images that we feed to our network.

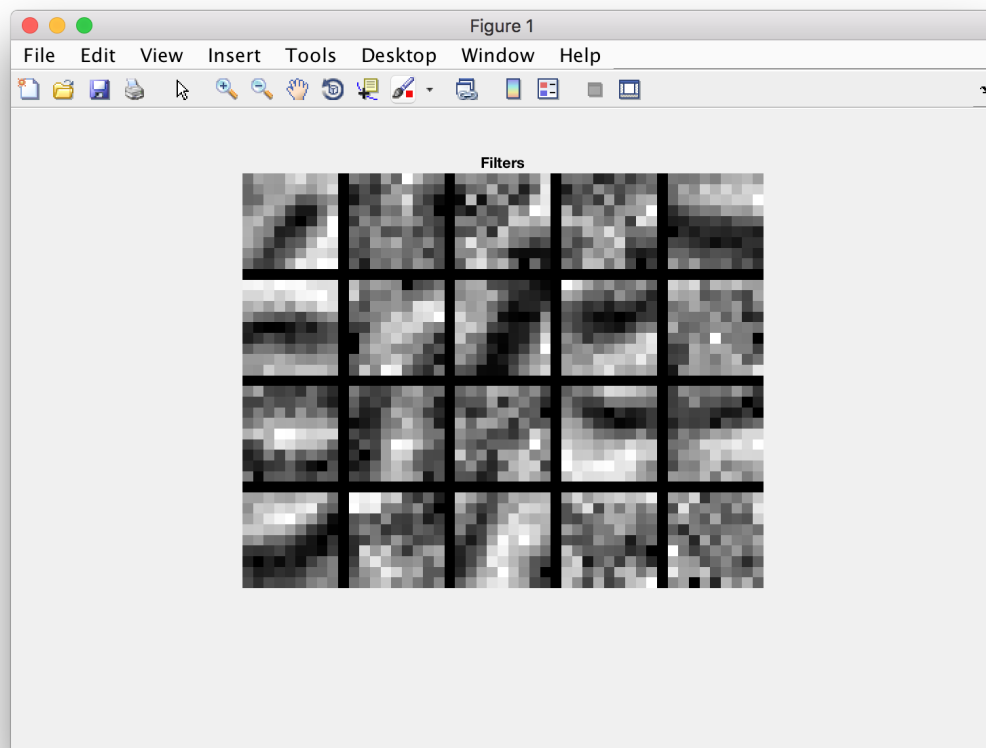


Figure 1: Filters after 3 epochs

3. Our network accuracy after 3 epochs is 97.24%